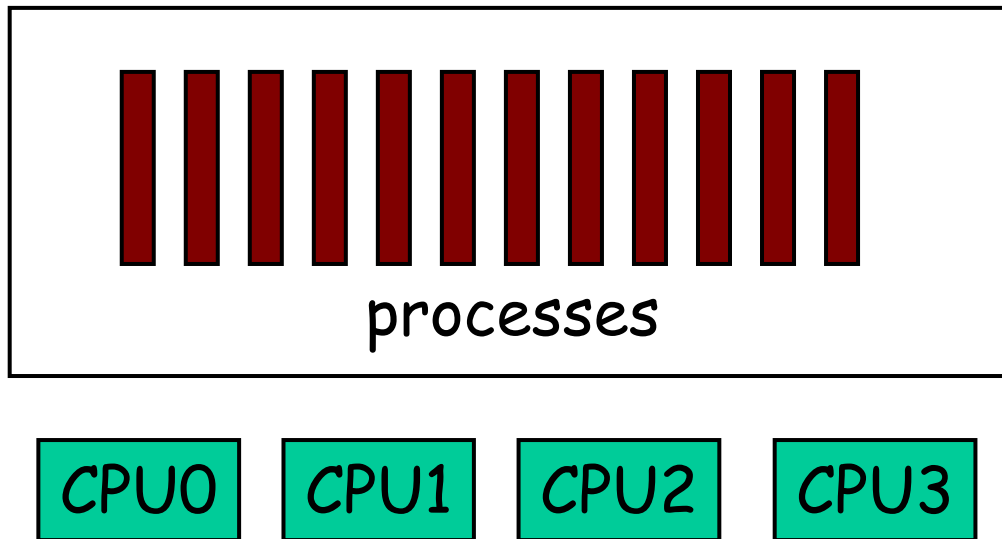


Multiprocessor scheduling issues


- ❑ Shared-memory Multiprocessor



- ❑ How to allocate processes to CPU?

Symmetric multiprocessor


- Architecture



- Small number of CPUs
- Same access time to main memory
- Private cache

Global queue of processes


- ❑ One ready queue shared across all CPUs



- ❑ Advantages
 - Good CPU utilization
 - Fair to all processes
- ❑ Disadvantages
 - Not scalable (contention for global queue lock)
 - Poor cache locality
- ❑ Linux 2.4 uses global queue

Per-CPU queue of processes


- ❑ Static partition of processes to CPUs



- ❑ Advantages
 - Easy to implement
 - Scalable (no contention on ready queue)
 - Better cache locality
- ❑ Disadvantages
 - Load-imbalance (some CPUs have more processes)
 - Unfair to processes and lower CPU utilization

Hybrid approach


- ❑ Use both global and per-CPU queues
- ❑ Balance jobs across queues



- ❑ Processor Affinity
 - Add process to a CPU's queue if recently run on the CPU
 - Cache state may still present
- ❑ Linux 2.6 uses a very similar approach

SMP: "gang" scheduling

- ❑ Multiple processes need coordination
- ❑ Should be scheduled simultaneously



- ❑ Scheduler on each CPU does not act independently
- ❑ **Coscheduling (gang scheduling)**: run a set of processes simultaneously
- ❑ **Global context-switch** across all CPUs