

```
!wget https://developer.nvidia.com/compute/cuda/9.2/Prod/local_installers/cuda-repo-ubuntu1604-9-2-local_9.2.88-1_amd64 -O cud
!dpkg -i cuda-repo-ubuntu1604-9-2-local_9.2.88-1_amd64.deb
!apt-key add /var/cuda-repo-9-2-local/7fa2af80.pub
!apt-get update
!apt-get install cuda-9.2
```

```
❏ --2023-01-21 03:26:14-- https://developer.nvidia.com/compute/cuda/9.2/Prod/local_installers/cuda-repo-ubuntu1604-9-2-lo
Resolving developer.nvidia.com (developer.nvidia.com)... 152.195.19.142
Connecting to developer.nvidia.com (developer.nvidia.com)|152.195.19.142|:443... connected.
HTTP request sent, awaiting response... 301 Moved Permanently
Location: https://developer.nvidia.com/downloads/compute/cuda/9.2/Prod/local_installers/cuda-repo-ubuntu1604-9-2-local_9.
--2023-01-21 03:26:15-- https://developer.nvidia.com/downloads/compute/cuda/9.2/Prod/local_installers/cuda-repo-ubuntu16
Reusing existing connection to developer.nvidia.com:443.
HTTP request sent, awaiting response... 302 Found
Location: https://developer.download.nvidia.com/compute/cuda/9.2/secure/Prod/local_installers/cuda-repo-ubuntu1604-9-2-lc
--2023-01-21 03:26:15-- https://developer.download.nvidia.com/compute/cuda/9.2/secure/Prod/local_installers/cuda-repo-u
Resolving developer.download.nvidia.com (developer.download.nvidia.com)... 152.195.19.142
Connecting to developer.download.nvidia.com (developer.download.nvidia.com)|152.195.19.142|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 1267391958 (1.2G) [application/x-deb]
Saving to: 'cuda-repo-ubuntu1604-9-2-local_9.2.88-1_amd64.deb'
```

```
cuda-repo-ubuntu160 100%[=====] 1.18G 197MB/s in 5.9s
```

```
2023-01-21 03:26:21 (204 MB/s) - 'cuda-repo-ubuntu1604-9-2-local_9.2.88-1_amd64.deb' saved [1267391958/1267391958]
```

```
Selecting previously unselected package cuda-repo-ubuntu1604-9-2-local.
(Reading database ... 129504 files and directories currently installed.)
Preparing to unpack cuda-repo-ubuntu1604-9-2-local_9.2.88-1_amd64.deb ...
Unpacking cuda-repo-ubuntu1604-9-2-local (9.2.88-1) ...
Setting up cuda-repo-ubuntu1604-9-2-local (9.2.88-1) ...
OK
Get:1 file:/var/cuda-repo-9-2-local InRelease
Ign:1 file:/var/cuda-repo-9-2-local InRelease
Get:2 file:/var/cuda-repo-9-2-local Release [574 B]
Get:2 file:/var/cuda-repo-9-2-local Release [574 B]
Get:3 file:/var/cuda-repo-9-2-local Release.gpg [819 B]
Get:3 file:/var/cuda-repo-9-2-local Release.gpg [819 B]
Get:4 https://cloud.r-project.org/bin/linux/ubuntu focal-cran40/ InRelease [3,622 B]
Hit:5 http://archive.ubuntu.com/ubuntu focal InRelease
Get:6 http://archive.ubuntu.com/ubuntu focal-updates InRelease [114 kB]
Ign:7 https://developer.download.nvidia.com/compute/machine-learning/repos/ubuntu2004/x86_64 InRelease
Hit:8 https://developer.download.nvidia.com/compute/cuda/repos/ubuntu2004/x86_64 InRelease
Hit:9 https://developer.download.nvidia.com/compute/machine-learning/repos/ubuntu2004/x86_64 Release
Get:10 http://security.ubuntu.com/ubuntu focal-security InRelease [114 kB]
Get:11 https://developer.download.nvidia.com/compute/machine-learning/repos/ubuntu2004/x86_64 InRelease [114 kB]
Hit:14 http://ppa.launchpad.net/cran/libgit2/ubuntu focal InRelease
Hit:15 http://ppa.launchpad.net/deadsnakes/ppa/ubuntu focal InRelease
Hit:16 http://ppa.launchpad.net/graphics-drivers/ppa/ubuntu focal InRelease
Get:18 http://archive.ubuntu.com/ubuntu focal-updates/universe amd64 Packages [1,284 kB]
Get:19 http://archive.ubuntu.com/ubuntu focal-updates/restricted amd64 Packages [2,003 kB]
Get:20 http://archive.ubuntu.com/ubuntu focal-updates/main amd64 Packages [2,909 kB]
Get:21 http://ppa.launchpad.net/c2d4u.team/c2d4u4.0+/ubuntu focal/main Sources [2,374 kB]
Get:22 http://security.ubuntu.com/ubuntu focal-security/universe amd64 Packages [982 kB]
Get:23 http://security.ubuntu.com/ubuntu focal-security/main amd64 Packages [2,436 kB]
Get:24 http://ppa.launchpad.net/c2d4u.team/c2d4u4.0+/ubuntu focal/main amd64 Packages [1,125 kB]
Fetched 13.5 MB in 4s (3,387 kB/s)
Reading package lists... Done
Reading package lists... Done
Building dependency tree
Reading state information... Done
Note, selecting 'cuda-9-2' for regex 'cuda-9.2'
```

Warning: You are connected to a GPU runtime, but not utilising the GPU. [Change to a standard runtime](#) X

```
!apt-get update
```

```
Get:1 file:/var/cuda-repo-9-2-local InRelease
Ign:1 file:/var/cuda-repo-9-2-local InRelease
Get:2 file:/var/cuda-repo-9-2-local Release [574 B]
Get:2 file:/var/cuda-repo-9-2-local Release [574 B]
Hit:3 https://cloud.r-project.org/bin/linux/ubuntu focal-cran40/ InRelease
Ign:4 https://developer.download.nvidia.com/compute/machine-learning/repos/ubuntu2004/x86_64 InRelease
Hit:5 https://developer.download.nvidia.com/compute/cuda/repos/ubuntu2004/x86_64 InRelease
Hit:6 https://developer.download.nvidia.com/compute/machine-learning/repos/ubuntu2004/x86_64 Release
Hit:7 http://security.ubuntu.com/ubuntu focal-security InRelease
Hit:8 http://ppa.launchpad.net/c2d4u.team/c2d4u4.0+/ubuntu focal InRelease
Hit:10 http://archive.ubuntu.com/ubuntu focal InRelease
Hit:11 http://archive.ubuntu.com/ubuntu focal-updates InRelease
Hit:12 http://ppa.launchpad.net/cran/libgit2/ubuntu focal InRelease
Hit:13 http://archive.ubuntu.com/ubuntu focal-backports InRelease
Hit:14 http://ppa.launchpad.net/deadsnakes/ppa/ubuntu focal InRelease
Hit:15 http://ppa.launchpad.net/graphics-drivers/ppa/ubuntu focal InRelease
Reading package lists... Done
```

```
!nvcc --version
```

```
nvcc: NVIDIA (R) Cuda compiler driver
Copyright (c) 2005-2018 NVIDIA Corporation
Built on Wed_Apr_11_23:16:29_CDT_2018
Cuda compilation tools, release 9.2, V9.2.88
```

```
!pip install git+https://github.com/andreinechaev/nvcc4jupyter.git
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
Collecting git+https://github.com/andreinechaev/nvcc4jupyter.git
  Cloning https://github.com/andreinechaev/nvcc4jupyter.git to /tmp/pip-req-build-9ydzhy1_
  Running command git clone --filter=blob:none --quiet https://github.com/andreinechaev/nvcc4jupyter.git /tmp/pip-req-bui
  Resolved https://github.com/andreinechaev/nvcc4jupyter.git to commit aac710a35f52bb78ab34d2e52517237941399eff
  Preparing metadata (setup.py) ... done
Building wheels for collected packages: NVCCPlugin
  Building wheel for NVCCPlugin (setup.py) ... done
  Created wheel for NVCCPlugin: filename=NVCCPlugin-0.0.2-py3-none-any.whl size=4304 sha256=41ea52abdf0ddcef598cce0fd835e
  Stored in directory: /tmp/pip-ephem-wheel-cache-1kf3321n/wheels/f3/08/cc/e2b5b0e1c92df07dbb50a6f024a68ce090f5e7b2316b41
Successfully built NVCCPlugin
Installing collected packages: NVCCPlugin
Successfully installed NVCCPlugin-0.0.2
```

```
%load_ext nvcc_plugin
```

```
created output directory at /content/src
Out bin /content/result.out
```

```
%%cu
#include <stdio.h>
__global__ void Hellokernel()
{
}
main()
{
Hellokernel <<<1, 1>>>();
printf("Hello cuda program srivani 22MAI1007\n");
return 0;
}
```

Warning: You are connected to a GPU runtime, but not utilising the GPU. [Change to a standard runtime](#) X

```
config.h:50,
runtime.h:78,
/usr/local/cuda/bin/../targets/x86_64-linux/include/crt/host_config.h:119:2: error: #error -- unsupported GNU version! gcc
119 | #error -- unsupported GNU version! gcc versions later than 7 are not supported!
    | ^~~~~
```

```
!pip install matplotlib-venn
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
Requirement already satisfied: matplotlib-venn in /usr/local/lib/python3.8/dist-packages (0.11.7)
Requirement already satisfied: scipy in /usr/local/lib/python3.8/dist-packages (from matplotlib-venn) (1.7.3)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.8/dist-packages (from matplotlib-venn) (3.2.2)
Requirement already satisfied: numpy in /usr/local/lib/python3.8/dist-packages (from matplotlib-venn) (1.21.6)
Requirement already satisfied: python-dateutil>=2.1 in /usr/local/lib/python3.8/dist-packages (from matplotlib->matplotlib)
Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in /usr/local/lib/python3.8/dist-packages (from matplotlib-venn)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.8/dist-packages (from matplotlib->matplotlib)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.8/dist-packages (from matplotlib->matplotlib-venn)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.8/dist-packages (from python-dateutil->matplotlib-venn)
```

```
%%cu
#include <stdio.h>
__global__ void add(int a, int b, int *c)
{
*c = a + b;
}
int main(void)
{
int c;
int *dev_c;
cudaMalloc((void**)&dev_c, sizeof(int));
add << <1, 1 >> > (2, 7, dev_c);
cudaMemcpy(&c, dev_c, sizeof(int),
cudaMemcpyDeviceToHost);
printf("2 + 7 = %d\n", c);
cudaFree(dev_c);
return 0;
}
```

```

In file included from /usr/local/cuda/bin/../targets/x86_64-linux/include/host_config.h:50,
                 from /usr/local/cuda/bin/../targets/x86_64-linux/include/cuda_runtime.h:78,
                 from <command-line>:
/usr/local/cuda/bin/../targets/x86_64-linux/include/crt/host_config.h:119:2: error: #error -- unsupported GNU version! gcc
119 | #error -- unsupported GNU version! gcc versions later than 7 are not supported!
    | ^~~~~

```

```

%%cu
#include <stdio.h>
__global__ void vector_add(int *out_d, int *a, int *b, int n)
{
    int bx = blockIdx.x;
    int by = blockIdx.y;
    int tx = threadIdx.x;
    int ty = threadIdx.y;
    int row = by*blockDim.y + ty;
    int col = bx*blockDim.x + tx;
    int dim = gridDim.x*blockDim.x;
    int i = row*dim + col;
    out_d[i] = a[i] + b[i];

}

int main()
{
    int *a, *b, *out_d,*out;
    int *d_a, *d_b;
    int N=6;
    int i;
    a = (int*)malloc(sizeof(int) * N);
    b = (int*)malloc(sizeof(int) * N);
    out = (int*)malloc(sizeof(int) * N);
    for (i=0;i<N;i++)
    {
        a[i]=i;
        b[i]=i*2;
    }
    cudaMalloc((void**)&d_a, sizeof(int) * N);
    cudaMalloc((void**)&d_b, sizeof(int) * N);
    cudaMalloc((void**)&out_d, sizeof(int) * N);
    cudaMemcpy(d_a, a, sizeof(int) * N, cudaMemcpyHostToDevice);
    cudaMemcpy(d_b, b, sizeof(int) * N, cudaMemcpyHostToDevice);

    printf( "Success ");
    for (i=0;i<N;i++)
    {
        printf("%d\n",out[i]);
    }
    cudaFree(d_a);
    cudaFree(d_b);
    cudaFree(out_d);
    free(a);
    free(b);
    free(out);
    return 0;
}

```

Warning: You are connected to a GPU runtime, but not utilising the GPU. [Change to a standard runtime](#) ✕

```

In file included from /usr/local/cuda/bin/../targets/x86_64-linux/include/host_config.h:50,
                 from /usr/local/cuda/bin/../targets/x86_64-linux/include/cuda_runtime.h:78,
                 from <command-line>:
/usr/local/cuda/bin/../targets/x86_64-linux/include/crt/host_config.h:119:2: error: #error -- unsupported GNU version! gcc
119 | #error -- unsupported GNU version! gcc versions later than 7 are not supported!
    | ^~~~~

```

```

%%cu
#include <stdio.h>
__global__ void matrixMul( int* Pd, int* Md, int* Nd, int width)
{
    int bx = blockIdx.x;
    int by = blockIdx.y;
    int tx = threadIdx.x;
    int ty = threadIdx.y;
    int col = by*blockDim.y + ty;
    int row = bx*blockDim.x + tx;
    int Pvalue=0;
    for (int k=0;k<width;++k)
        Pvalue+=Md[row*width+k]*Nd[k*width+col];
    Pd[row*width+col]=Pvalue;
}

```

```

}

int main()
{
    int *M, *N1, *Md, *Nd, *Pd, *P;
    const int xb = 3; /* gridDim.x */
    const int yb = 3; /* gridDim.y */
    const int zb = 1; /* gridDim.z */
    const int xt = 3; /* blockDim.x */
    const int yt = 3; /* blockDim.y */
    const int zt = 1; /* blockDim.z */
    int N,width;
    int i;
    width=9;
    N=width*width;
    M = (int*)malloc(sizeof(int) * N);
    N1 = (int*)malloc(sizeof(int) * N);
    P = (int*)malloc(sizeof(int) * N);
    for (i=0;i<N;i++)
    {
        M[i]=i;
        N1[i]=i*2;
    }
    dim3 dimGrid(xb,yb,zb);
    dim3 dimBlock(xt,yt,zt);
    cudaMalloc((void**)&Md, sizeof(int) * N);
    cudaMalloc((void**)&Nd, sizeof(int) * N);
    cudaMalloc((void**)&Pd, sizeof(int) * N);
    cudaMemcpy(Md, M, sizeof(int) * N, cudaMemcpyHostToDevice);
    cudaMemcpy(Nd, N1, sizeof(int) * N, cudaMemcpyHostToDevice);
    matrixMul<<<dimGrid,dimBlock>>>(Pd, Md, Nd, width);
    cudaMemcpy(P, Pd, sizeof(int) * N, cudaMemcpyDeviceToHost);
    printf("Success");
    for (i=0;i<N;i++)
    {
        printf("%d\n",P[i]);
    }
    cudaFree(Md);
    cudaFree(Nd);
    cudaFree(Pd);
    free(M);
    free(N1);
    free(P);
}

```

Warning: You are connected to a GPU runtime, but not utilising the GPU. [Change to a standard runtime](#) ✕

```

In file included from /usr/local/cuda/bin/../targets/x86_64-linux/include/host_config.h:50,
                 from /usr/local/cuda/bin/../targets/x86_64-linux/include/cuda_runtime.h:78,
                 from <command-line>:
/usr/local/cuda/bin/../targets/x86_64-linux/include/crt/host_config.h:119:2: error: #error -- unsupported GNU version! gcc
119 | #error -- unsupported GNU version! gcc versions later than 7 are not supported!
    | ^~~~~

```

```

%%cu
#include <stdio.h>
__global__ void matrixFill ( int *x )
{
    int bx = blockDim.x;
    int by = blockDim.y;
    int tx = threadIdx.x;
    int ty = threadIdx.y;
    int col = by*blockDim.y + ty;
    int row = bx*blockDim.x + tx;
    int dim =blockDim.x*gridDim.x;
    int i = row*dim + col;
    x[i] = i;
}
int main ( int argc, char* argv[] )
{
    const int xb = 2; /* gridDim.x */
    const int yb = 2; /* gridDim.y */
    const int zb = 1; /* gridDim.z */
    const int xt = 2; /* blockDim.x */
    const int yt = 2; /* blockDim.y */
    const int zt = 1; /* blockDim.z */
    const int n = xb*yb*zb*xt*yt*zt;
    printf("allocating array of length %d...\n",n);
    int *xhost = (int*)calloc(n,sizeof(int));
    for(int i=0; i<n; i++) xhost[i] = -1.0;
    int *xdevice;
}

```

```

size_t sx = n*sizeof(int);
cudaMalloc((void**)&xdevice,sx);
cudaMemcpy(xdevice,xhost,sx,cudaMemcpyHostToDevice);
dim3 dimGrid(xb,yb,zb);
dim3 dimBlock(xt,yt,zt);
matrixFill<<<dimGrid,dimBlock>>>(xdevice);
cudaMemcpy(xhost,xdevice,sx,cudaMemcpyDeviceToHost);
cudaFree(xdevice);
int *p = xhost;
for(int i1=0; i1 < xb; i1++)
for(int i2=0; i2 < yb; i2++)
for(int i3=0; i3 < zb; i3++)
for(int i4=0; i4 < xt; i4++)
for(int i5=0; i5 < yt; i5++)
for(int i6=0; i6 < zt; i6++)
printf("x[%d][%d][%d][%d][%d] = %d\n",i1,i2,i3,i4,i5,i6,*(p++));
return 0;
}

```

```

In file included from /usr/local/cuda/bin/../targets/x86_64-linux/include/host_config.h:50,
                 from /usr/local/cuda/bin/../targets/x86_64-linux/include/cuda_runtime.h:78,
                 from <command-line>:
/usr/local/cuda/bin/../targets/x86_64-linux/include/crt/host_config.h:119:2: error: #error -- unsupported GNU version! gcc
119 | #error -- unsupported GNU version! gcc versions later than 7 are not supported!
    | ^~~~~

```

Warning: You are connected to a GPU runtime, but not utilising the GPU. [Change to a standard runtime](#) ✕

[Colab paid products](#) - [Cancel contracts here](#)

✓ 0s completed at 08:58

