

```
!apt-get update
```

```
Get:1 http://security.ubuntu.com/ubuntu bionic-security InRelease [88.7 kB]
Get:2 https://cloud.r-project.org/bin/linux/ubuntu bionic-cran40/ InRelease [3,626 B]
Ign:3 https://developer.download.nvidia.com/compute/machine-learning/repos/ubuntu1804/x86\_64 InRelease
Hit:4 https://developer.download.nvidia.com/compute/cuda/repos/ubuntu1804/x86\_64 InRelease
Hit:5 https://developer.download.nvidia.com/compute/machine-learning/repos/ubuntu1804/x86\_64 Release
Get:6 http://ppa.launchpad.net/c2d4u.team/c2d4u4.0+/ubuntu bionic InRelease [15.9 kB]
Hit:7 http://archive.ubuntu.com/ubuntu bionic InRelease
Get:8 http://archive.ubuntu.com/ubuntu bionic-updates InRelease [88.7 kB]
Hit:9 http://ppa.launchpad.net/cran/libgit2/ubuntu bionic InRelease
Hit:10 http://ppa.launchpad.net/deadsnakes/ppa/ubuntu bionic InRelease
Get:11 http://security.ubuntu.com/ubuntu bionic-security/main amd64 Packages [3,127 kB]
Get:12 http://archive.ubuntu.com/ubuntu bionic-backports InRelease [83.3 kB]
Get:13 http://ppa.launchpad.net/graphics-drivers/ppa/ubuntu bionic InRelease [21.3 kB]
Get:14 http://security.ubuntu.com/ubuntu bionic-security/restricted amd64 Packages [1,349 kB]
Get:15 http://security.ubuntu.com/ubuntu bionic-security/universe amd64 Packages [1,573 kB]
Get:17 http://ppa.launchpad.net/c2d4u.team/c2d4u4.0+/ubuntu bionic/main Sources [2,240 kB]
Get:18 http://archive.ubuntu.com/ubuntu bionic-updates/main amd64 Packages [3,552 kB]
Get:19 http://ppa.launchpad.net/c2d4u.team/c2d4u4.0+/ubuntu bionic/main amd64 Packages [1,145 kB]
Get:20 http://archive.ubuntu.com/ubuntu bionic-updates/restricted amd64 Packages [1,392 kB]
Get:21 http://archive.ubuntu.com/ubuntu bionic-updates/multiverse amd64 Packages [30.8 kB]
Get:22 http://archive.ubuntu.com/ubuntu bionic-updates/universe amd64 Packages [2,348 kB]
Get:23 http://ppa.launchpad.net/graphics-drivers/ppa/ubuntu bionic/main amd64 Packages [43.2 kB]
Fetched 17.1 MB in 4s (4,430 kB/s)
Reading package lists... Done
```

```
!wget https://developer.nvidia.com/compute/cuda/9.2/Prod/local_installers/cuda-repo-ubuntu1604-9-2-local_9.2.88-1_amd64 -O cuda-repo-ubuntu1604-9-2-local_9.2.88-1_amd64.deb
!dpkg -i cuda-repo-ubuntu1604-9-2-local_9.2.88-1_amd64.deb
!apt-key add /var/cuda-repo-9-2-local/7fa2af80.pub
!apt-get update
!apt-get install cuda-9.2
```

```
➤ --2023-01-10 13:19:32-- https://developer.nvidia.com/compute/cuda/9.2/Prod/local\_installers/cuda-repo-ubuntu1604-9-2-local\_9.2.88-1\_amd64.deb
Resolving developer.nvidia.com (developer.nvidia.com)... 152.195.19.142
Connecting to developer.nvidia.com (developer.nvidia.com)|152.195.19.142|:443... connected.
HTTP request sent, awaiting response... 301 Moved Permanently
Location: https://developer.nvidia.com/downloads/compute/cuda/9.2/Prod/local\_installers/cuda-repo-ubuntu1604-9-2-local\_9.2.88-1\_amd64.deb
--2023-01-10 13:19:32-- https://developer.nvidia.com/downloads/compute/cuda/9.2/Prod/local\_installers/cuda-repo-ubuntu1604-9-2-local\_9.2.88-1\_amd64.deb
Reusing existing connection to developer.nvidia.com:443.
HTTP request sent, awaiting response... 302 Found
Location: https://developer.download.nvidia.com/compute/cuda/9.2/secure/Prod/local\_installers/cuda-repo-ubuntu1604-9-2-local\_9.2.88-1\_amd64.deb
--2023-01-10 13:19:32-- https://developer.download.nvidia.com/compute/cuda/9.2/secure/Prod/local\_installers/cuda-repo-ubuntu1604-9-2-local\_9.2.88-1\_amd64.deb
Resolving developer.download.nvidia.com (developer.download.nvidia.com)... 152.195.19.142
Connecting to developer.download.nvidia.com (developer.download.nvidia.com)|152.195.19.142|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 1267391958 (1.2G) [application/x-deb]
Saving to: 'cuda-repo-ubuntu1604-9-2-local_9.2.88-1_amd64.deb'
```

```
cuda-repo-ubuntu160 100%[=====] 1.18G 169MB/s in 6.7s
```

```
2023-01-10 13:19:39 (181 MB/s) - 'cuda-repo-ubuntu1604-9-2-local_9.2.88-1_amd64.deb' saved [1267391958/1267391958]
```

```
Selecting previously unselected package cuda-repo-ubuntu1604-9-2-local.
(Reading database ... 124016 files and directories currently installed.)
Preparing to unpack cuda-repo-ubuntu1604-9-2-local_9.2.88-1_amd64.deb ...
Unpacking cuda-repo-ubuntu1604-9-2-local (9.2.88-1) ...
Setting up cuda-repo-ubuntu1604-9-2-local (9.2.88-1) ...
OK
Get:1 file:/var/cuda-repo-9-2-local InRelease
Ign:1 file:/var/cuda-repo-9-2-local InRelease
Get:2 file:/var/cuda-repo-9-2-local Release [574 B]
Get:2 file:/var/cuda-repo-9-2-local Release [574 B]
Get:3 file:/var/cuda-repo-9-2-local Release.gpg [819 B]
Get:3 file:/var/cuda-repo-9-2-local Release.gpg [819 B]
Hit:4 https://cloud.r-project.org/bin/linux/ubuntu bionic-cran40/ InRelease
Ign:5 https://developer.download.nvidia.com/compute/machine-learning/repos/ubuntu1804/x86\_64 InRelease
Hit:6 https://developer.download.nvidia.com/compute/cuda/repos/ubuntu1804/x86\_64 InRelease
Hit:7 http://archive.ubuntu.com/ubuntu bionic InRelease
Hit:8 http://security.ubuntu.com/ubuntu bionic-security InRelease
Hit:9 http://ppa.launchpad.net/c2d4u.team/c2d4u4.0+/ubuntu bionic InRelease
Hit:10 https://developer.download.nvidia.com/compute/machine-learning/repos/ubuntu1804/x86\_64 Release
Get:11 file:/var/cuda-repo-9-2-local Packages [18.7 kB]
Hit:12 http://archive.ubuntu.com/ubuntu bionic-updates InRelease
Hit:13 http://archive.ubuntu.com/ubuntu bionic-backports InRelease
Hit:14 http://ppa.launchpad.net/cran/libgit2/ubuntu bionic InRelease
Hit:15 http://ppa.launchpad.net/deadsnakes/ppa/ubuntu bionic InRelease
Hit:16 http://ppa.launchpad.net/graphics-drivers/ppa/ubuntu bionic InRelease
Reading package lists... Done
Reading package lists... Done
Building dependency tree
Reading state information... Done
Note, selecting 'cuda-9-2' for regex 'cuda-9.2'
Note, selecting 'libcuda-9.2-1' for regex 'cuda-9.2'
The following package was automatically installed and is no longer required:
  libnvidia-common-460
Use 'apt autoremove' to remove it.
```

The following additional packages will be installed:  
cuda-command-line-tools-9-2 cuda-compiler-9-2 cuda-cublas-9-2  
cuda-cublas-dev-9-2 cuda-cudart-9-2 cuda-cudart-dev-9-2 cuda-cufft-9-2  
cuda-cufft-dev-9-2 cuda-cuobidump-9-2 cuda-cuopti-9-2 cuda-curand-9-2

```
!nvcc --version
```

```
nvcc: NVIDIA (R) Cuda compiler driver
Copyright (c) 2005-2018 NVIDIA Corporation
Built on Wed_Apr_11_23:16:29_CDT_2018
Cuda compilation tools, release 9.2, V9.2.88
```

```
!pip install git+https://github.com/andreinechaev/nvcc4jupyter.git
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
Collecting git+https://github.com/andreinechaev/nvcc4jupyter.git
  Cloning https://github.com/andreinechaev/nvcc4jupyter.git to /tmp/pip-req-build-5y92v3oc
  Running command git clone --filter=blob:none --quiet https://github.com/andreinechaev/nvcc4jupyter.git /tmp/pip-req-bui
  Resolved https://github.com/andreinechaev/nvcc4jupyter.git to commit aac710a35f52bb78ab34d2e52517237941399eff
  Preparing metadata (setup.py) ... done
Building wheels for collected packages: NVCCPlugin
  Building wheel for NVCCPlugin (setup.py) ... done
  Created wheel for NVCCPlugin: filename=NVCCPlugin-0.0.2-py3-none-any.whl size=4304 sha256=354b497e248d08b83a2aa3cc0238(
  Stored in directory: /tmp/pip-ephem-wheel-cache-xguwhln8/wheels/f3/08/cc/e2b5b0e1c92df07dbb50a6f024a68ce090f5e7b2316b41
Successfully built NVCCPlugin
Installing collected packages: NVCCPlugin
Successfully installed NVCCPlugin-0.0.2
```

```
%load_ext nvcc_plugin
```

```
created output directory at /content/src
Out bin /content/result.out
```

%%cu - extension for cuda programming

\_\_ global \_\_ is a functional qualiifier called from the **CPU** and lead to the **GPU**

HelloKernel() user defined function

Hellokernel<<<1, 1>>>(): The number of cuda threads that execute that kernel for a given kernel call. No. of cuda threads going to executed.

<<<1, 1>>>: <<<block size, thread no>>>

```
%%cu
#include <stdio.h>
__global__ void Hellokernel()
{
}
main()
{
Hellokernel<<<1, 1>>>();
printf("Hello World\n");
return 0;
}
```

```
Hello World
```

```
%%cu
#include <stdio.h>
__global__ void add(int a, int b, int *c)
{
*c = a + b;
}
int main(void)
{
int c;
int *dev_c;
cudaMalloc((void**)&dev_c, sizeof(int));
add << <1, 1 >> > (2, 7, dev_c);
cudaMemcpy(&c, dev_c, sizeof(int),
cudaMemcpyDeviceToHost);
printf("2 + 7 = %d\n", c);
cudaFree(dev_c);
return 0;
}
```

```
2 + 7 = 1
```

```

%%cu
#include <stdio.h>
__global__ void vector_add(int *out_d, int *a, int *b, int n)
{
    int bx = blockIdx.x;
int by = blockIdx.y;
int tx = threadIdx.x;
int ty = threadIdx.y;
int row = by*blockDim.y + ty;
int col = bx*blockDim.x + tx;
int dim = gridDim.x*blockDim.x;
int i = row*dim + col;
out_d[i] = a[i] + b[i];

}

int main()
{
    int *a, *b, *out_d,*out;
    int *d_a, *d_b;
    int N=6;
    int i;
    a = (int*)malloc(sizeof(int) * N);
    b = (int*)malloc(sizeof(int) * N);
    out = (int*)malloc(sizeof(int) * N);
    for (i=0;i<N;i++)
    {
        a[i]=i;
        b[i]=i*2;
    }
    cudaMalloc((void**)&d_a, sizeof(int) * N);
    cudaMalloc((void**)&d_b, sizeof(int) * N);
    cudaMalloc((void**)&out_d, sizeof(int) * N);
    cudaMemcpy(d_a, a, sizeof(int) * N, cudaMemcpyHostToDevice);
    cudaMemcpy(d_b, b, sizeof(int) * N, cudaMemcpyHostToDevice);
    vector_add<<<2,4>>>(out_d, d_a, d_b, N);
    cudaMemcpy(out, out_d, sizeof(int) * N, cudaMemcpyDeviceToHost);
    printf("Success");
    for (i=0;i<N;i++)
    {
        printf("%d\n",out[i]);
    }
    cudaFree(d_a);
    cudaFree(d_b);
    cudaFree(out_d);
    free(a);
    free(b);
    free(out);
    return 0;
}

```

```

Success1760174240
22085
0
0
0
0
0

```

```

%%cu
#include <stdio.h>
__global__ void matrixMul( int* Pd, int* Md, int* Nd, int width)
{
    int bx = blockIdx.x;
int by = blockIdx.y;
int tx = threadIdx.x;
int ty = threadIdx.y;
int col = by*blockDim.y + ty;
int row = bx*blockDim.x + tx;
int Pvalue=0;
for (int k=0;k<width;++k)
    Pvalue+=Md[row*width+k]*Nd[k*width+col];
Pd[row*width+col]=Pvalue;
}

int main()
{
    int *M, *N1, *Md, *Nd, *Pd, *P;
    const int xb = 3; /* gridDim.x */
    const int yb = 3; /* gridDim.y */
    const int zb = 1; /* gridDim.z */

```



0  
0  
0  
0  
0  
0  
0  
0  
0

```
%%cu
#include <stdio.h>
__global__ void matrixFill ( int *x )
{
    int bx = blockIdx.x;
    int by = blockIdx.y;
    int tx = threadIdx.x;
    int ty = threadIdx.y;
    int col = by*blockDim.y + ty;
    int row = bx*blockDim.x + tx;
    int dim =blockDim.x*gridDim.x;
    int i = row*dim + col;
    x[i] = i;
}
int main ( int argc, char* argv[] )
{
    const int xb = 2; /* gridDim.x */
    const int yb = 2; /* gridDim.y */
    const int zb = 1; /* gridDim.z */
    const int xt = 2; /* blockDim.x */
    const int yt = 2; /* blockDim.y */
    const int zt = 1; /* blockDim.z */
    const int n = xb*yb*zb*xt*yt*zt;
    printf("allocating array of length %d...\n",n);
    int *xhost = (int*)calloc(n,sizeof(int));
    for(int i=0; i<n; i++) xhost[i] = -1.0;
    int *xdevice;
    size_t sx = n*sizeof(int);
    cudaMalloc((void**)&xdevice,sx);
    cudaMemcpy(xdevice,xhost,sx,cudaMemcpyHostToDevice);
    dim3 dimGrid(xb,yb,zb);
    dim3 dimBlock(xt,yt,zt);
    matrixFill<<<dimGrid,dimBlock>>>(xdevice);
    cudaMemcpy(xhost,xdevice,sx,cudaMemcpyDeviceToHost);
    cudaFree(xdevice);
    int *p = xhost;
    for(int i1=0; i1 < xb; i1++)
    for(int i2=0; i2 < yb; i2++)
    for(int i3=0; i3 < zb; i3++)
    for(int i4=0; i4 < xt; i4++)
    for(int i5=0; i5 < yt; i5++)
    for(int i6=0; i6 < zt; i6++)
    printf("x[%d][%d][%d][%d][%d][%d] = %d\n",i1,i2,i3,i4,i5,i6,*(p++));
    return 0;
}
```

```
allocating array of length 16...
x[0][0][0][0][0][0] = -1
x[0][0][0][0][1][0] = -1
x[0][0][0][1][0][0] = -1
x[0][0][0][1][1][0] = -1
x[0][1][0][0][0][0] = -1
x[0][1][0][0][1][0] = -1
x[0][1][0][1][0][0] = -1
x[0][1][0][1][1][0] = -1
x[1][0][0][0][0][0] = -1
x[1][0][0][0][1][0] = -1
x[1][0][0][1][0][0] = -1
x[1][0][0][1][1][0] = -1
x[1][1][0][0][0][0] = -1
x[1][1][0][0][1][0] = -1
x[1][1][0][1][0][0] = -1
x[1][1][0][1][1][0] = -1
```

---

✓ 0s completed at 18:52

